

389I: Healthcare & AI

ter.ps/389iweek1



Topics

- Syllabus Week + Intro to ML Algorithms
- Machine Learning for Healthcare
- Early Disease Detection
- Electronic Health Records (EHRs)
- Drug Discovery
- Telemedicine
- Mental Health




Learning Outcomes

- Understand the central role of data in improving healthcare
- Understand healthcare data-specific ethics and challenges
- Be able to apply basic ML techniques to analyze healthcare datasets



Assignments

- Lectures
 - Overview of relevant AI algorithms & domain specific case studies
- Quizzes
 - In class on lecture content or assigned readings
- Codelabs
 - Guided projects on analyzing specific datasets to achieve set results
- Final project
 - Chose a dataset in a relevant domain to analyze & present findings to class



Codelabs

- We will provide datasets in a specific domain to analyze
- Given starter code with directions in Jupyter notebook
- Apply ML techniques & write analysis
- 2 this semester + final project



Instructors

- **Max Leiserson**
 - mdml@cs.umd.edu
 - [website](#)
- **Sanna Madan**
 - smadan12@umd.edu
- **Kyle Liu**
 - kliu1234@umd.edu
- **Office Hours**
 - TBD / by appointment
- **Piazza**



Machine Learning

What is machine learning?

- “Labeling machine.” Teaching a computer how to **classify**.
 - complex, hard-to-see patterns
- More specifically:

training of a model from data
that generalizes a decision
against a performance measure.



What is “artificial intelligence”?

- Much broader than machine learning
- Computers mimicking the cognitive functions of humans
 - intelligence
- We're not quite there yet

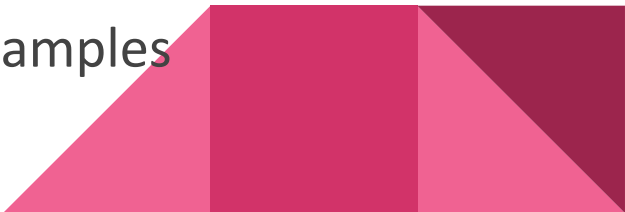


Machine learning

- Focuses on learning from datasets provided to the computer
- Identify patterns and make decisions from these datasets
- Self driving cars, Netflix recommendations, fraud detection, etc.



Deep learning

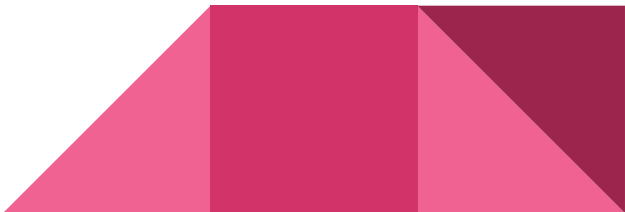
- Even *deeper* - a subset of machine learning
 - Multi-layer machine learning model with hidden layers
 - Blackbox
 - is this ideal for healthcare?
 - Needs much larger datasets, often millions of samples
- 

In this course, we are focusing on **machine learning**.
Not deep learning, not AI.

Why?



Supervised learning

- Most practical machine learning is **supervised**
 - Input variables and output variables are given: $Y = F(x)$
 - Learns mapping function from input \rightarrow output
 - Labels are **known** beforehand
 - Classification, regression
- 

Unsupervised learning

- Only input data, with no corresponding output labels: $? = F(x)$
- No correct answers to learn from
- Algorithms find their own “structures” in the data instead
- Clustering, association



Data

- Machine learning doesn't matter if the data is of poor quality!!!!!!!
- If the data is flawed, then the machine learning will be flawed
- Cleansing data may even take up most of the “machine learning” process
- To trust the insights of the data, you must first trust the data



Quality datasets are therefore the first step to ML.



Data cont'd

- Amount of data is also crucial
- More is more
- Large datasets in healthcare/bio once rare, now up and coming



Data cont'd

- Bias in data must also be considered
- e.g., people who get their DNA sequenced tend to be wealthier
- Machine may be getting a skewed perception of reality



Summary

- Machine learning is only a subset of artificial intelligence
- There are many different approaches to machine learning
- Quality, quantity, and bias in data are crucial considerations



Links

- [Github](#)
- [Survey](#)
- Slides: ter.ps/389iweek1
- [Tips for ML](#)

